

**Deccan Education Society's (DES)
Pune University, Pune
School of Engineering and Technology
Department of Computer Engineering and Technology
Program: B. Tech in Computer Science and Engineering**

Academic Year: 2024-25	Year: Second Year	Semester: II
PRN No.: 1012412079	Name: Ratnajeet Patil	
Subject: Database Management System		
Assignment No.: 4		
Date:		

Lab Assignment: 04

Title: Advanced SQL Queries: Design and execute following SQL queries on suitable applications.

Joins: Retrieve data from multiple related tables using INNER JOIN, LEFT JOIN, and RIGHT JOIN.

Aggregation: Use aggregate functions like COUNT, AVG, SUM, MIN, MAX in combination with GROUP BY and HAVING clauses.

Theory:

What is JOIN in SQL?

A JOIN in SQL is used to combine rows from two or more tables based on a related column between them. It allows retrieving data from multiple tables in a single query.

Different Types of JOINS

1. INNER JOIN – Returns only the matching rows between the two tables.

```
1. SELECT A.*, B.*  
2. FROM TableA A  
3. INNER JOIN TableB B  
4. ON A.common_column = B.common_column;  
5.
```

2. LEFT JOIN (LEFT OUTER JOIN) – Returns all rows from the left table and the matched rows from the right table. If no match is found, NULLs are returned.

```
1. SELECT A.*, B.*  
2. FROM TableA A  
3. LEFT JOIN TableB B  
4. ON A.common_column = B.common_column;  
5.
```

3. RIGHT JOIN (RIGHT OUTER JOIN) – Returns all rows from the right table and the matched rows from the left table.

**Deccan Education Society's (DES)
Pune University, Pune
School of Engineering and Technology
Department of Computer Engineering and Technology
Program: B. Tech in Computer Science and Engineering**

```
1. SELECT A.*, B.*  
2. FROM TableA A  
3. RIGHT JOIN TableB B  
4. ON A.common_column = B.common_column;  
5.
```

4. FULL JOIN (FULL OUTER JOIN) – Returns all rows when there is a match in either table. If there's no match, NULLs are returned.

```
1. SELECT A.*, B.*  
2. FROM TableA A  
3. FULL JOIN TableB B  
4. ON A.common_column = B.common_column;  
5.
```

5. CROSS JOIN – Returns the Cartesian product of both tables (every row in TableA is combined with every row in TableB).

```
1. SELECT A.*, B.*  
2. FROM TableA A  
3. CROSS JOIN TableB B;  
4.
```

6. SELF JOIN – A table joins itself by using aliases.

```
1. SELECT A.*, B.*  
2. FROM TableA A  
3. INNER JOIN TableA B  
4. ON A.column_name = B.column_name;  
5.
```

Aggregate Functions

Aggregate functions perform calculations on a set of values and return a single value.

Function Description

COUNT() Returns the number of rows

SUM() Returns the total sum of a column

AVG() Returns the average value of a column

MAX() Returns the highest value in a column

MIN() Returns the lowest value in a column

**Deccan Education Society's (DES)
Pune University, Pune
School of Engineering and Technology
Department of Computer Engineering and Technology
Program: B. Tech in Computer Science and Engineering**

Group By

The GROUP BY clause is used with aggregate functions to group rows that have the same values in specified columns.

Example:

```
1. SELECT department, COUNT(*) AS total_employees
2. FROM employees
3. GROUP BY department;
4.
```

Having

The HAVING clause is used to filter groups created by GROUP BY based on aggregate function results (since WHERE cannot be used with aggregate functions).

Show Query Execution Screenshots for:

- Inner join
- Left join
- Right join
- Cross join
- Self join
- Full Outer join
- All Aggregate Functions
- Group By
- Having
- Set Operators (Union, Intersect and Difference)
- Use of Date, Time, Mathematical and String Functions

**Deccan Education Society's (DES)
Pune University, Pune
School of Engineering and Technology**

```

-- Create sample tables
CREATE TABLE customers (
  customer_id SERIAL PRIMARY KEY,
  name VARCHAR(100),
  email VARCHAR(100),
  city VARCHAR(50),
  join_date DATE
);

CREATE TABLE orders (
  order_id SERIAL PRIMARY KEY,
  customer_id INT,
  order_date DATE,
  total_amount DECIMAL(10,2)
);

CREATE TABLE products (
  product_id SERIAL PRIMARY KEY,
  name VARCHAR(100),
  category VARCHAR(50),
  price DECIMAL(10,2)
);

CREATE TABLE order_items (
  item_id SERIAL PRIMARY KEY,
  order_id INT,
  product_id INT,
  quantity INT,
  price DECIMAL(10,2)
);

-- Insert sample data
INSERT INTO customers (name, email, city, join_date) VALUES
('John Smith', 'john@example.com', 'New York', '2023-01-15'),
('Mary Johnson', 'mary@example.com', 'Los Angeles', '2023-02-10'),
('Robert Brown', 'robert@example.com', 'Chicago', '2023-03-05'),
('Patricia Davis', 'patricia@example.com', 'Houston', '2023-04-20'),
('Michael Miller', 'michael@example.com', 'Phoenix', '2023-05-12');

INSERT INTO orders (customer_id, order_date, total_amount) VALUES
(1, '2023-06-10', 120.50),
(2, '2023-06-15', 85.25),
(1, '2023-07-01', 200.00),
(3, '2023-07-05', 65.75),
(2, '2023-07-10', 150.30),
(NULL, '2023-07-15', 50.00);

INSERT INTO products (name, category, price) VALUES
('Laptop', 'Electronics', 1200.00),
('Smartphone', 'Electronics', 800.00),
('Headphones', 'Electronics', 150.00),
('Coffee Maker', 'Kitchen', 85.00),
('Blender', 'Kitchen', 65.00);

INSERT INTO order_items (order_id, product_id, quantity, price) VALUES
(1, 1, 1, 1200.00),
(1, 3, 1, 150.00),
(2, 2, 1, 800.00),
(3, 4, 2, 65.00),
(3, 5, 1, 65.00),
(4, 3, 1, 150.00),
(5, 4, 1, 85.00);

-- SQL Shell (pgsql)
CREATE TABLE employees (
  employee_id SERIAL PRIMARY KEY,
  name VARCHAR(100),
  manager_id INT
);

CREATE TABLE employees (name, manager_id) VALUES
('John Smith', NULL),
('Mary Johnson', 1),
('Robert Brown', 1),
('Patricia Davis', 2),
('Michael Miller', 2);

-- Inner Join
SELECT c.name, o.order_id, o.order_date, o.total_amount
FROM customers c
INNER JOIN orders o ON c.customer_id = o.customer_id;

-- Left Join
SELECT c.name, o.order_id, o.order_date, o.total_amount
FROM customers c
LEFT JOIN orders o ON c.customer_id = o.customer_id;

-- Right Join
SELECT c.name, o.order_id, o.order_date, o.total_amount
FROM customers c
RIGHT JOIN orders o ON c.customer_id = o.customer_id;

```

**Deccan Education Society's (DES)
Pune University, Pune
School of Engineering and Technology
Department of Computer Engineering and Technology
Program: B. Tech in Computer Science and Engineering**

```
SQL Shell (psql)
name | order_id | order_date | total_amount
-----|-----|-----|-----
John Smith | 1 | 2023-06-10 | 120.50
Mary Johnson | 2 | 2023-06-15 | 85.25
John Smith | 3 | 2023-07-01 | 200.00
Robert Brown | 4 | 2023-07-05 | 65.75
Mary Johnson | 5 | 2023-07-10 | 150.30
Mary Johnson | 6 | 2023-07-15 | 50.00
(6 rows)

sbaspac=#
sbaspac=# -- Cross Join
sbaspac=# SELECT c.name, p.name AS product_name
sbaspac=# FROM customers c
sbaspac=# CROSS JOIN products p
sbaspac=# LIMIT 10;
name | product_name
-----|-----
John Smith | Laptop
Mary Johnson | Laptop
Robert Brown | Laptop
Patricia Davis | Laptop
Michael Miller | Laptop
John Smith | Smartphone
Mary Johnson | Smartphone
Robert Brown | Smartphone
Patricia Davis | Smartphone
Michael Miller | Smartphone
(10 rows)

sbaspac=#
sbaspac=# -- Self Join
sbaspac=# SELECT e.name AS employee, m.name AS manager
sbaspac=# FROM employees e
sbaspac=# LEFT JOIN employees m ON e.manager_id = m.employee_id;
employee | manager
-----|-----
John Smith | John Smith
Mary Johnson | John Smith
Patricia Davis | Mary Johnson
Michael Miller | Mary Johnson
(5 rows)

sbaspac=#
sbaspac=# -- Full Outer Join
sbaspac=# SELECT c.name, o.order_id, o.order_date, o.total_amount
sbaspac=# FROM customers c
sbaspac=# FULL OUTER JOIN orders o ON c.customer_id = o.customer_id;
name | order_id | order_date | total_amount
-----|-----|-----|-----
John Smith | 1 | 2023-06-10 | 120.50
Mary Johnson | 2 | 2023-06-15 | 85.25
John Smith | 3 | 2023-07-01 | 200.00
Robert Brown | 4 | 2023-07-05 | 65.75
Mary Johnson | 5 | 2023-07-10 | 150.30
Mary Johnson | 6 | 2023-07-15 | 50.00
Michael Miller |
Patricia Davis |
(8 rows)

SQL Shell (psql)

sbaspac=#
sbaspac=# -- All Aggregate Functions
sbaspac=# SELECT
sbaspac=# COUNT(*) AS total_orders,
sbaspac=# SUM(total_amount) AS total_sales,
sbaspac=# AVG(total_amount) AS average_order_value,
sbaspac=# MIN(total_amount) AS min_order_value,
sbaspac=# MAX(total_amount) AS max_order_value,
sbaspac=# STDDEV(total_amount) AS standard_deviation,
sbaspac=# VARIANCE(total_amount) AS variance
sbaspac=# FROM orders;
total_orders | total_sales | average_order_value | min_order_value | max_order_value | standard_deviation | variance
-----|-----|-----|-----|-----|-----|-----
(1 row)
6 | 671.80 | 111.96666666666667 | 50.00 | 200.00 | 56.536286289410353 | 3196.3516666666667

sbaspac=#
sbaspac=# -- Group By
sbaspac=# SELECT customer_id, COUNT(*) AS order_count, SUM(total_amount) AS total_spent
sbaspac=# FROM orders
sbaspac=# WHERE customer_id IS NOT NULL
sbaspac=# GROUP BY customer_id;
customer_id | order_count | total_spent
-----|-----|-----
3 | 1 | 65.75
2 | 2 | 235.55
1 | 2 | 320.50
(3 rows)

sbaspac=#
sbaspac=# -- Having
sbaspac=# SELECT customer_id, COUNT(*) AS order_count, SUM(total_amount) AS total_spent
sbaspac=# FROM orders
sbaspac=# WHERE customer_id IS NOT NULL
sbaspac=# GROUP BY customer_id
sbaspac=# HAVING COUNT(*) > 1;
customer_id | order_count | total_spent
-----|-----|-----
2 | 2 | 235.55
1 | 2 | 320.50
(2 rows)

sbaspac=#
sbaspac=# -- Set Operators - UNION
sbaspac=# SELECT name FROM customers
sbaspac=# UNION
sbaspac=# SELECT name FROM employees;
name
-----
Michael Miller
Mary Johnson
John Smith
Robert Brown
Patricia Davis
(6 rows)

sbaspac=#
```

Deccan Education Society's (DES)
Pune University, Pune
School of Engineering and Technology
Department of Computer Engineering and Technology
Program: B. Tech in Computer Science and Engineering

```
SQL Shell (psql)

dbmsprac=# -- Set Operators - UNION ALL
dbmsprac=# SELECT name FROM customers
dbmsprac=# UNION ALL
dbmsprac=# SELECT name FROM employees;
name
-----
John Smith
Mary Johnson
Robert Brown
Patricia Davis
Michael Miller
John Smith
Mary Johnson
Robert Brown
Patricia Davis
Michael Miller
(10 rows)

dbmsprac=#
dbmsprac=# -- Set Operators - INTERSECT
dbmsprac=# SELECT city FROM customers
dbmsprac=# INTERSECT
dbmsprac=# SELECT name FROM products;
city
-----
(0 rows)

dbmsprac=#
dbmsprac=# -- Set Operators - EXCEPT
dbmsprac=# SELECT name FROM customers
dbmsprac=# EXCEPT
dbmsprac=# SELECT name FROM employees;
name
-----
(0 rows)

dbmsprac=#
dbmsprac=# -- Date functions
dbmsprac=# SELECT
dbmsprac=#     order_date,
dbmsprac=#     EXTRACT(YEAR FROM order_date) AS year,
dbmsprac=#     EXTRACT(MONTH FROM order_date) AS month,
dbmsprac=#     EXTRACT(DAY FROM order_date) AS day,
dbmsprac=#     CURRENT_DATE AS today,
dbmsprac=#     order_date + INTERVAL '30 days' AS due_date,
dbmsprac=#     AGE(CURRENT_DATE, order_date) AS order_age
dbmsprac=# FROM orders;
 order_date | year | month | day | today          | due_date          | order_age
-----
2023-06-10 | 2023 | 6     | 10  | 2025-03-17    | 2023-07-10 00:00:00 | 1 year 9 mons 7 days
2023-06-15 | 2023 | 6     | 15  | 2025-03-17    | 2023-07-15 00:00:00 | 1 year 9 mons 2 days
2023-07-01 | 2023 | 7     | 1   | 2025-03-17    | 2023-07-31 00:00:00 | 1 year 8 mons 16 days
2023-07-05 | 2023 | 7     | 5   | 2025-03-17    | 2023-08-04 00:00:00 | 1 year 8 mons 12 days
2023-07-10 | 2023 | 7     | 10  | 2025-03-17    | 2023-08-09 00:00:00 | 1 year 8 mons 7 days
2023-07-15 | 2023 | 7     | 15  | 2025-03-17    | 2023-08-14 00:00:00 | 1 year 8 mons 2 days
(6 rows)

SQL Shell (psql)

dbmsprac=#
dbmsprac=# -- Time functions
dbmsprac=# SELECT
dbmsprac=#     CURRENT_TIME AS current_time,
dbmsprac=#     CURRENT_TIMESTAMP AS current_timestamp,
dbmsprac=#     LOCALTIME AS local_time,
dbmsprac=#     LOCALTIMESTAMP AS local_timestamp,
dbmsprac=#     EXTRACT(HOUR FROM CURRENT_TIME) AS current_hour;
 current_time | current_timestamp | local_time | local_timestamp | current_hour
-----
19:13:55.287464+05:30 | 2025-03-17 19:13:55.287464+05:30 | 19:13:55.287464 | 2025-03-17 19:13:55.287464 | 19
(1 row)

dbmsprac=#
dbmsprac=# -- Mathematical functions
dbmsprac=# SELECT
dbmsprac=#     price,
dbmsprac=#     ROUND(price, 1) AS rounded_price,
dbmsprac=#     CEIL(price) AS ceiling_price,
dbmsprac=#     FLOOR(price) AS floor_price,
dbmsprac=#     ABS(price) AS absolute_price,
dbmsprac=#     POWER(price, 2) AS price_squared,
dbmsprac=#     SQRT(price) AS square_root,
dbmsprac=#     LOG(price) AS natural_log,
dbmsprac=#     RANDOM() AS random_number
dbmsprac=# FROM products;
 price | rounded_price | ceiling_price | floor_price | absolute_price | price_squared | square_root | natural_log | random_number
-----
1200.00 | 1200.0       | 1200         | 1200        | 1200.00        | 1440000.0000000000 | 34.641016151377546 | 3.0791812468476248 | 0.07243349673741428
800.00 | 800.0        | 800          | 800         | 800.00         | 640000.000000000000 | 28.284271247461901 | 2.96580979869919436 | 0.10039526011060724
150.00 | 150.0        | 150          | 150         | 150.00         | 22500.000000000000 | 12.247481719915800 | 2.1768932598536812 | 0.0511375544672118
85.00 | 85.0         | 85           | 85          | 85.00          | 7225.00000000000000 | 9.21944457292887 | 1.9294189257142927 | 0.6739132384765105
65.00 | 65.0         | 65           | 65          | 65.00          | 4225.00000000000000 | 8.062257748298550 | 1.8129133566428550 | 0.44349828254968544
(5 rows)

dbmsprac=#
dbmsprac=# -- String functions
dbmsprac=# SELECT
dbmsprac=#     name,
dbmsprac=#     UPPER(name) AS upper_case,
dbmsprac=#     LOWER(name) AS lower_case,
dbmsprac=#     LENGTH(name) AS name_length,
dbmsprac=#     SUBSTRING(name, 1, 3) AS first_three_chars,
dbmsprac=#     REPLACE(name, 'a', 'A') AS replaced,
dbmsprac=#     CONCAT(name, ' ', category) AS product_info,
dbmsprac=#     TRIM(BOTH ' ' FROM name) AS trimmed_name
dbmsprac=# FROM products;
 name | upper_case | lower_case | name_length | first_three_chars | replaced | product_info | trimmed_name
-----
Laptop | LAPTOP     | laptop     | 6           | Lap               | Laptop   | Laptop - Electronics | Laptop
Smartphone | SMARTPHONE | smartphone | 10          | Sma               | Smartphone | Smartphone - Electronics | Smartphone
Headphones | HEADPHONES | headphones | 10          | Hea               | Headphones | Headphones - Electronics | Headphones
Coffee Maker | COFFEE MAKER | coffee maker | 12          | Cof               | Coffee Maker | Coffee Maker - Kitchen | Coffee Maker
Blender | BLENDER   | blender    | 7           | Ble               | Blender   | Blender - Kitchen | Blender
(5 rows)

dbmsprac=#
```

**Deccan Education Society's (DES)
Pune University, Pune
School of Engineering and Technology
Department of Computer Engineering and Technology
Program: B. Tech in Computer Science and Engineering**

FAQs:

1. Can aggregate functions be used without the GROUP BY clause?

Yes, aggregate functions can be used without GROUP BY. When used without GROUP BY, they apply to the entire result set and return a single row. For example:

```
1. SELECT COUNT(*), AVG(total_amount), SUM(total_amount) FROM orders;
```

2. What is the difference between COUNT(*) and COUNT(column_name)?

COUNT(*) counts all rows in the table, including rows with NULL values

COUNT(column_name) counts only non-NULL values in the specified column. For example, if a column has NULL values, COUNT(column_name) will return a smaller number than COUNT(*).

3. Can we use multiple aggregate functions in a single query? Yes, you can use multiple aggregate functions in a single query. For example:

```
1. SELECT COUNT(*) AS total_orders,  
2.    SUM(total_amount) AS total_sales,  
3.    AVG(total_amount) AS average_order  
4. FROM orders;  
5.
```

4. Can we use WHERE and HAVING together in a query with aggregate functions? Yes, you can use both WHERE and HAVING together. The key difference is:

WHERE filters rows before they're grouped

HAVING filters groups after they're formed. For example:

```
1. SELECT customer_id, COUNT(*) AS order_count  
2. FROM orders  
3. WHERE order_date > '2023-01-01'  
4. GROUP BY customer_id  
5. HAVING COUNT(*) > 5;  
6.
```

5. What is the difference between INNER JOIN and LEFT JOIN?

INNER JOIN: Returns only rows that have matching values in both tables

**Deccan Education Society's (DES)
Pune University, Pune
School of Engineering and Technology
Department of Computer Engineering and Technology
Program: B. Tech in Computer Science and Engineering**

LEFT JOIN: Returns all rows from the left table and matching rows from the right table. If no match is found, NULL values are returned for right table columns

6. What is the difference between ON and USING in joins?

ON: Allows you to specify any join condition (with any columns)

```
1. -- Using ON
2. SELECT * FROM orders o JOIN customers c ON o.customer_id = c.customer_id;
3.
4.
```

USING: Simplifies the syntax when joining tables on columns with the same name For example

```
-- 1. Using USING (works only when column names are identical)
2. SELECT * FROM orders JOIN customers USING (customer_id);
3.
```

7. Can we use aggregate functions with joins?

Yes, you can use aggregate functions with joins. This is often used to aggregate data from related tables. For example:

```
1. SELECT c.name, COUNT(o.order_id) AS order_count, SUM(o.total_amount) AS total_spent
2. FROM customers c
3. LEFT JOIN orders o ON c.customer_id = o.customer_id
4. GROUP BY c.name;
5.
```

This query joins customers and orders tables, then aggregates the order data for each customer.